100% success.
One missing day.

Both rates are correct. Only one begins before the missing record.

DELIVERY VIEW

100%

29 delivered

29 recorded

One absent day unseen

SCHEDULE VIEW

93.3%

28 accepted

30 expected

Every day retained

Illustration only. No live reliability rate or OpenAI failure is claimed.

> **A perfect delivery score can hide an incomplete schedule.**

Green begins too late.

Five evidence states answer five different questions.



States 1 and 5 belong to the application contract; the middle states are platform evidence.

> **The first record must exist before the invocation does.**

The denominator changes the truth.

Start with the promise, not the records that survived long enough to be counted.

Accepted task outputs

Expected scheduled occurrences

Transport → 29 delivered / 29 recorded = 100%

Outcome → 28 accepted / 30 expected = 93.3%

Status → Illustration; not observed performance

> **Outcome reliability begins with expected occurrences.**

One missing record changes every layer.

The same illustrative month, shown as counts rather than percentages.

OCCURRENCES OR OUTPUTS



Illustrative counts only. Causes remain unassigned.

> **A narrowing funnel is evidence; an absent first row is a blind spot.**

Keep three scoreboards.

Transport, execution and outcome should not borrow one another's labels.

LAYER	NUMERATOR	DENOMINATOR
Transport	Delivered invocations	Recorded invocations
Execution	Terminal root turns	Delivered invocations
Outcome	Accepted artifacts	Expected occurrences

Split execution outcomes into completed, failed and cancelled; inspect tools and output.

> **Name the layer before choosing the fix.**

Build one row before work begins.

A stable occurrence key lets missing evidence stay missing instead of disappearing.

01

1 · Expect

Create schedule ID plus scheduled_at before the run.

02

2 · Join

Attach invocation, session and root-turn identities.

03

3 · Inspect

Record tool results, artifact hash and acceptance.

04

4 · Recover

Keep the occurrence key; give retries new attempt IDs.

Measure timing from intended occurrence to accepted artifact when lateness matters.

> **Join attempts forward; never rebuild the promise from surviving attempts.**

Who counted the days?

A system cannot prove completeness using only records created after the missing point.

THE EASY SCORE

Success among records found

Accepted work among runs promised

Build the expected-occurrence ledger before the next scheduled wake-up.

Illustrative counts · No live test · No platform failure claim

> **Reliability earns trust at the denominator, not the green badge.**